

Міністерство освіти і науки України
Національний технічний університет України «Київський політехнічний
інститут імені Ігоря Сікорського»
Факультет інформатики та обчислювальної техніки

Кафедра інформатики та програмної інженерії

Звіт
з лабораторної роботи № 4 з дисципліни
«Програмування інтелектуальних інформаційних систем»

Виконав студент ІІІ-12 Басараб Олег Андрійович

Перевірів Баришич Лука Маріянович

Київ 2023

Lab-4

```
import warnings
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import VotingClassifier
from sklearn.datasets import load_wine
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.linear_model import RidgeClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import AdaBoostClassifier,
GradientBoostingClassifier
from xgboost import XGBClassifier
from sklearn.linear_model import Ridge, Lasso, LogisticRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.base import BaseEstimator, RegressorMixin, clone

warnings.filterwarnings(action='ignore')
seed = 42
```

Import data and data preprocessing

```
df = pd.read_csv("resources/data.csv")
df.drop(['Unnamed: 32', 'id'], axis=1, inplace=True)
df['diagnosis'] = df['diagnosis'].astype('category').cat.codes
X = df.drop(['diagnosis'], axis = 1)
y = df['diagnosis']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = seed)

df.head()
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	1	17.99	10.38	122.80	1001.0	
1	1	20.57	17.77	132.90	1326.0	
2	1	19.69	21.25	130.00	1203.0	

3	1	11.42	20.38	77.58	386.1
4	1	20.29	14.34	135.10	1297.0

	smoothness_mean	compactness_mean	concavity_mean	concave
points_mean \				
0	0.11840	0.27760	0.3001	
0.14710				
1	0.08474	0.07864	0.0869	
0.07017				
2	0.10960	0.15990	0.1974	
0.12790				
3	0.14250	0.28390	0.2414	
0.10520				
4	0.10030	0.13280	0.1980	
0.10430				

	symmetry_mean	...	radius_worst	texture_worst	perimeter_worst	\
0	0.2419	...	25.38	17.33	184.60	
1	0.1812	...	24.99	23.41	158.80	
2	0.2069	...	23.57	25.53	152.50	
3	0.2597	...	14.91	26.50	98.87	
4	0.1809	...	22.54	16.67	152.20	

	area_worst	smoothness_worst	compactness_worst	concavity_worst	\
0	2019.0	0.1622	0.6656	0.7119	
1	1956.0	0.1238	0.1866	0.2416	
2	1709.0	0.1444	0.4245	0.4504	
3	567.7	0.2098	0.8663	0.6869	
4	1575.0	0.1374	0.2050	0.4000	

	concave	points_worst	symmetry_worst	fractal_dimension_worst
0		0.2654	0.4601	0.11890
1		0.1860	0.2750	0.08902
2		0.2430	0.3613	0.08758
3		0.2575	0.6638	0.17300
4		0.1625	0.2364	0.07678

[5 rows x 31 columns]

Hyperparameter tuning

Decision Tree

```
base_dt = DecisionTreeClassifier()
base_dt.fit(X_train, y_train)
base_dt_y_pred = base_dt.predict(X_test)

parameters = {'max_features': ['log2', 'sqrt', 'auto'],
              'criterion': ['entropy', 'gini'],
```

```

        'max_depth': [2, 3, 5, 10, 50],
        'min_samples_split': [2, 3, 50, 100],
        'min_samples_leaf': [1, 5, 8, 10]
    }

    grid_obj = GridSearchCV(base_dt, parameters)
    grid_obj = grid_obj.fit(X_train, y_train)

    tuned_dt = grid_obj.best_estimator_
    tuned_dt.fit(X_train, y_train)
    tuned_dt_y_pred = tuned_dt.predict(X_test)

    acc_base_dt = round(metrics.accuracy_score(y_test, base_dt_y_pred) *
100, 2)
    acc_tuned_dt = round(metrics.accuracy_score(y_test, tuned_dt_y_pred) *
100, 2)

    print('Accuracy of base Decision Tree model: ', acc_base_dt)
    print('Accuracy of tuned Decision Tree model: ', acc_tuned_dt)

    Accuracy of base Decision Tree model: 94.15
    Accuracy of tuned Decision Tree model: 97.08

```

Random Forest

```

base_rf = RandomForestClassifier()
base_rf.fit(X_train, y_train)
base_rf_y_pred = base_rf.predict(X_test)

parameters = {'n_estimators': [4, 6, 9, 10, 15],
        'max_features': ['log2', 'sqrt', 'auto'],
        'criterion': ['entropy', 'gini'],
        'max_depth': [2, 3, 5, 10],
        'min_samples_split': [2, 3, 5],
        'min_samples_leaf': [1, 5, 8]
    }

    grid_obj = GridSearchCV(base_rf, parameters)
    grid_obj = grid_obj.fit(X_train, y_train)

    tuned_rf = grid_obj.best_estimator_
    tuned_rf.fit(X_train, y_train)
    tuned_rf_y_pred = tuned_rf.predict(X_test)

    acc_base_rf = round(metrics.accuracy_score(y_test, base_rf_y_pred) *
100, 2)
    acc_tuned_rf = round(metrics.accuracy_score(y_test, tuned_rf_y_pred) *
100, 2)

```

```
print('Accuracy of base Random Forest model: ', acc_base_rf)
print('Accuracy of tuned Random Forest model: ', acc_tuned_rf)
```

```
Accuracy of base Random Forest model: 96.49
Accuracy of tuned Random Forest model: 97.08
```

Support Vector Machine

```
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

base_svc = SVC()
base_svc.fit(X_train, y_train)
base_svc_y_pred = base_svc.predict(X_test)

parameters = [
    {'C': [1, 10, 100, 1000], 'kernel': ['linear']},
    {'C': [1, 10, 100, 1000], 'gamma': [0.001, 0.0001], 'kernel':
 ['rbf']},
]

grid_obj = GridSearchCV(base_svc, parameters)
grid_obj = grid_obj.fit(X_train, y_train)

tuned_svc = grid_obj.best_estimator_
tuned_svc.fit(X_train, y_train)
tuned_svc_y_pred = tuned_svc.predict(X_test)

acc_base_svc = round(metrics.accuracy_score(y_test, base_svc_y_pred) *
100, 2)
acc_tuned_svc = round(metrics.accuracy_score(y_test, tuned_svc_y_pred)
* 100, 2)

print('Accuracy of base SVC model: ', acc_base_svc)
print('Accuracy of tuned SVC model: ', acc_tuned_svc)

Accuracy of base SVC model: 97.66
Accuracy of tuned SVC model: 97.66
```

K-Nearest Neighbors

```
base_knn = KNeighborsClassifier()
base_knn.fit(X_train, y_train)
base_knn_y_pred = base_knn.predict(X_test)

parameters = {'n_neighbors': [3, 4, 5, 10],
              'weights': ['uniform', 'distance'],
              'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],
              'leaf_size': [10, 20, 30, 50]}
```

```

    }

grid_obj = GridSearchCV(base_knn, parameters)
grid_obj = grid_obj.fit(X_train, y_train)

tuned_knn = grid_obj.best_estimator_
tuned_knn.fit(X_train, y_train)
tuned_knn_y_pred = tuned_knn.predict(X_test)

acc_base_knn = round(metrics.accuracy_score(y_test, base_knn_y_pred) *
100, 2)
acc_tuned_knn = round(metrics.accuracy_score(y_test, tuned_knn_y_pred)
* 100, 2)

print('Accuracy of base KNN model: ', acc_base_knn)
print('Accuracy of tuned KNN model: ', acc_tuned_knn)

Accuracy of base KNN model: 95.91
Accuracy of tuned KNN model: 95.91

```

Tuning results

```

models = pd.DataFrame({
    'Model': ['Base Decision Tree', 'Tuned Decision Tree', 'Base
Random Forest', 'Tuned Random Forest',
            'Base Support Vector Machines', 'Tuned Support Vector
Machines', 'Base K-Nearest Neighbors', 'Tuned K-Nearest Neighbors'],
    'Accuracy': [acc_base_dt, acc_tuned_dt, acc_base_rf, acc_tuned_rf,
                acc_base_svc, acc_tuned_svc, acc_base_knn,
acc_tuned_knn]})

models.sort_values(by='Accuracy', ascending=False)

```

	Model	Accuracy
4	Base Support Vector Machines	97.66
5	Tuned Support Vector Machines	97.66
1	Tuned Decision Tree	97.08
3	Tuned Random Forest	97.08
2	Base Random Forest	96.49
6	Base K-Nearest Neighbors	95.91
7	Tuned K-Nearest Neighbors	95.91
0	Base Decision Tree	94.15

Max Voting

```

estimators = []
estimators.append(('LR', LogisticRegression(solver='lbfgs',
multi_class='multinomial', max_iter=200)))
estimators.append(('SVC', SVC(gamma='auto', probability=True)))

```

```

estimators.append(('DTC', DecisionTreeClassifier()))

hard_voting = VotingClassifier(estimators=estimators, voting='hard')
hard_voting.fit(X_train, y_train)
y_pred = hard_voting.predict(X_test)

score = metrics.accuracy_score(y_test, y_pred)
print("Accuracy of Hard Voting model: %f" % score)

soft_voting = VotingClassifier(estimators=estimators, voting='soft')
soft_voting.fit(X_train, y_train)
y_pred = soft_voting.predict(X_test)

score = metrics.accuracy_score(y_test, y_pred)
print("Accuracy of Soft Voting model: %f" % score)

Accuracy of Hard Voting model: 0.988304
Accuracy of Soft Voting model: 0.988304

```

Weighted Averaging

```

class AverageWeight(BaseEstimator, RegressorMixin):
    def __init__(self, model, weight):
        self.model = model
        self.weight = weight

    def fit(self, X, y):
        self.models_ = [clone(x) for x in self.model]
        for model in self.models_:
            model.fit(X, y)
        return self

    def predict(self, X):
        w = list()
        pred = np.array([model.predict(X) for model in self.models_])
        # for every data point, single model prediction times weight,
then add them together
        for data in range(pred.shape[1]):
            single = [pred[model, data]*weight for model, weight in
zip(range(pred.shape[0]), self.weight)]
            w.append(np.sum(single))
        return w

    def rmse_cv(model, X, y):
        rmse = np.sqrt(-
cross_val_score(model, X, y, scoring="neg_mean_squared_error", cv=5))
        return rmse

estimators = []
estimators.append(LogisticRegression())

```

```

estimators.append(DecisionTreeRegressor())
estimators.append(Lasso())
estimators.append(Ridge())

w1 = 0.2
w2 = 0.3
w3 = 0.4
w4 = 0.1

weight_avg = AverageWeight(model=estimators, weight=[w1, w2, w3, w4])
score = rmse_cv(weight_avg, X, y)

print("Accuracy of Weighted Averaging model: %f" % score.mean())

Accuracy of Weighted Averaging model: 0.234249

```

Blending

```

X_train, X_val, y_train, y_val = train_test_split(X_train, y_train,
test_size=0.25, random_state=seed)

x_val = pd.DataFrame(X_val)
x_test = pd.DataFrame(X_test)

model1 = DecisionTreeClassifier()
model1.fit(X_train, y_train)
val_pred1=model1.predict(X_val)
test_pred1=model1.predict(X_test)
val_pred1=pd.DataFrame(val_pred1)
test_pred1=pd.DataFrame(test_pred1)

model2 = KNeighborsClassifier()
model2.fit(X_train,y_train)
val_pred2 = model2.predict(X_val)
test_pred2 = model2.predict(X_test)
val_pred2 = pd.DataFrame(val_pred2)
test_pred2 = pd.DataFrame(test_pred2)

df_val = pd.concat([x_val, val_pred1, val_pred2], axis=1)
df_test = pd.concat([x_test, test_pred1, test_pred2], axis=1)

model = LogisticRegression()
model.fit(df_val, y_val)
print("Accuracy of Blending model: ", model.score(df_test, y_test))

Accuracy of Blending model: 0.9941520467836257

```


Bagging

```
rf = RandomForestClassifier()
et = ExtraTreesClassifier()
knn = KNeighborsClassifier()
svc = SVC()
rg = RidgeClassifier()
clf_array = [rf, et, knn, svc, rg]

for clf in clf_array:
    vanilla_scores = cross_val_score(clf, X, y, cv=10, n_jobs=-1)
    bagging_clf = BaggingClassifier(clf, max_samples=0.4,
    max_features=10, random_state=seed)
    bagging_scores = cross_val_score(bagging_clf, X, y, cv=10, n_jobs=-1)

    print ("Mean of: {1:.3f}, std: (+/-) {2:.3f} [{0}]".format(clf.__class__.__name__, vanilla_scores.mean(),
    vanilla_scores.std()))
    print ("Mean of: {1:.3f}, std: (+/-) {2:.3f} [Bagging {0}]\n".format(clf.__class__.__name__, bagging_scores.mean(),
    bagging_scores.std()))

Mean of: 0.961, std: (+/-) 0.030 [RandomForestClassifier]
Mean of: 0.954, std: (+/-) 0.037 [Bagging RandomForestClassifier]

Mean of: 0.967, std: (+/-) 0.028 [ExtraTreesClassifier]
Mean of: 0.953, std: (+/-) 0.032 [Bagging ExtraTreesClassifier]

Mean of: 0.930, std: (+/-) 0.029 [KNeighborsClassifier]
Mean of: 0.931, std: (+/-) 0.029 [Bagging KNeighborsClassifier]

Mean of: 0.914, std: (+/-) 0.029 [SVC]
Mean of: 0.910, std: (+/-) 0.042 [Bagging SVC]

Mean of: 0.954, std: (+/-) 0.025 [RidgeClassifier]
Mean of: 0.935, std: (+/-) 0.024 [Bagging RidgeClassifier]

clf = [rf, et, knn, svc, rg]
eclf = VotingClassifier(estimators=[('Random Forests', rf), ('Extra Trees', et), ('KNeighbors', knn), ('SVC', svc), ('Ridge Classifier', rg)], voting='hard')
for clf, label in zip([rf, et, knn, svc, rg, eclf], ['Random Forest', 'Extra Trees', 'KNeighbors', 'SVC', 'Ridge Classifier', 'Ensemble']):
    scores = cross_val_score(clf, X, y, cv=10, scoring='accuracy')
    print("Accuracy: %0.2f (+/- %0.2f) [%s]" % (scores.mean(), scores.std(), label))

Accuracy: 0.96 (+/- 0.03) [Random Forest]
Accuracy: 0.96 (+/- 0.02) [Extra Trees]
```

```
Accuracy: 0.93 (+/- 0.03) [KNeighbors]
Accuracy: 0.91 (+/- 0.03) [SVC]
Accuracy: 0.95 (+/- 0.03) [Ridge Classifier]
Accuracy: 0.96 (+/- 0.02) [Ensemble]
```

Boosting

```
ada_boost = AdaBoostClassifier(random_state=seed)
ada_boost.fit(X_train, y_train)
ada_boost.score(X_test, y_test)

grad_boost =
GradientBoostingClassifier(learning_rate=0.01, random_state=seed)
grad_boost.fit(X_train, y_train)
grad_boost.score(X_test, y_test)

xgb_boost = XGBClassifier(random_state=1, learning_rate=0.01)
xgb_boost.fit(X_train, y_train)
xgb_boost.score(X_test, y_test)

eclf = VotingClassifier(estimators=[('Ada Boost', ada_boost), ('Grad
Boost', grad_boost), ('XG Boost', xgb_boost)], voting='hard')
clf = [rf, et, knn, svc, rg]
for clf, label in zip([ada_boost, grad_boost, xgb_boost, eclf], ['Ada
Boost', 'Grad Boost', 'XG Boost', 'Ensemble']):
    scores = cross_val_score(clf, X, y, cv=10, scoring='accuracy')
    print("Accuracy: %0.2f (+/- %0.2f) [%s]" % (scores.mean(),
    scores.std(), label))

Accuracy: 0.96 (+/- 0.03) [Ada Boost]
Accuracy: 0.95 (+/- 0.02) [Grad Boost]
Accuracy: 0.95 (+/- 0.02) [XG Boost]
Accuracy: 0.96 (+/- 0.02) [Ensemble]
```